



VLM bandwidth: explicit DRAM scenario and measured L2 traffic

If every tensor read/write reached DRAM



Actual L2 byte counter / NCU kernel time



Left: single-read/write tensor model; cache hits, rereads and deferred writeback are unknown. Right: L2 is not DRAM.